

Bios 6301 Assignment 5

Rui Wang

11/10/2016

Due Tuesday, 15 November, 1:00 PM

$5^{n=\text{day}}$ points taken off for each day late.

50 points total.

Submit a single knitr file (named `homework5.rmd`), along with a valid PDF output file. Inside the file, clearly indicate which parts of your responses go with which problems (you may use the original homework document as a template). Add your name as `author` to the file's metadata section. Raw R code/output or word processor files are not acceptable.

Failure to name file `homework5.rmd` or include author name may result in 5 points taken off.

Question 1

24 points

Import the HAART dataset (`haart.csv`) from the GitHub repository into R, and perform the following manipulations: (4 points each)

```
url <- "https://github.com/fonnesbeck/Bios6301/raw/master/datasets/haart.csv"
haart <- read.csv(url, header = TRUE, stringsAsFactors = FALSE)
haart$init.date <- as.Date(haart$init.date, "%m/%d/%y")
haart$last.visit <- as.Date(haart$last.visit, "%m/%d/%y")
haart$date.death <- as.Date(haart$date.death, "%m/%d/%y")
```

1. Convert date columns into a usable (for analysis) format. Use the `table` command to display the counts of the year from `init.date`.

```
table(format(haart$init.date, "%Y"))
```

```
##
## 1998 2000 2001 2002 2003 2004 2005 2006 2007
##      1      5     17     60    270   292   207   104    44
```

2. Create an indicator variable (one which takes the values 0 or 1 only) to represent death within 1 year of the initial visit. How many observations died in year 1?

```
# eliminate more than 1 year of the initial visit and marked NA in the date.death
haart$indicator <- ifelse((haart$date.death - haart$init.date > 365 | is.na(haart$date.death)), 0, 1)
sum(haart$indicator==1)
```

```
## [1] 92
```

3. Use the `init.date`, `last.visit` and `death.date` columns to calculate a followup time (in days), which is the difference between the first and either the last visit or a death event (whichever comes first). If these times are longer than 1 year, censor them (this means if the value is above 365, set followup to 365). Print the quantile for this new variable.

```
# if NA in the last.visit, use the date.death - init
# if
haart$follow.up <- ifelse(is.na(haart$last.visit), haart$date.death - haart$init.date,
                        haart$last.visit - haart$init.date)
```

```
haart$follow.up[haart$follow.up > 365] <- 365
quantile(haart$follow.up)
```

```
##      0%      25%      50%      75%     100%
##    0.00 320.75 365.00 365.00 365.00
```

4. Create another indicator variable representing loss to followup; this means the observation is not known to be dead but does not have any followup visits after the first year. How many records are lost-to-followup?

```
haart$loss.to.followup <- ifelse(haart$follow.up < 365 & haart$death == 0, 1,0)
sum(haart$loss.to.followup)
```

```
## [1] 173
```

5. Recall our work in class, which separated the `init.reg` field into a set of indicator variables, one for each unique drug. Create these fields and append them to the database as new columns. Which drug regimen are found over 100 times?

```
# create a list for all users
row.reg <- strsplit(haart[, 'init.reg'], ",")
# find all the unique drugs
all.reg <- unique(unlist(strsplit(haart[, 'init.reg'], ",")))
# create a database for user.reg
user.reg <- sapply(all.reg, function(j) sapply(row.reg, function(i) j %in% i))
# append it to haart database
haart <- cbind(haart, user.reg)
# check which drug regimen are found over 100 times
colnames(user.reg)[colSums(user.reg) > 100]
```

```
## [1] "3TC" "AZT" "EFV" "NVP" "D4T"
```

6. The dataset `haart2.csv` contains a few additional observations for the same study. Import these and append them to your master dataset (if you were smart about how you coded the previous steps, cleaning the additional observations should be easy!). Show the first five records and the last five records of the complete (and clean) data set.

```
url1 <- "https://github.com/fonnesbeck/Bios6301/raw/master/datasets/haart.csv"
haart <- read.csv(url1, header = TRUE, stringsAsFactors = FALSE)
url2 <- "https://github.com/fonnesbeck/Bios6301/raw/master/datasets/haart2.csv"
haart2 <- read.csv(url2, header = TRUE, stringsAsFactors = FALSE)
# combine these two databases
haart <- rbind(haart, haart2)
# copy the code from 1-5
haart$init.date <- as.Date(haart$init.date, "%m/%d/%y")
haart$last.visit <- as.Date(haart$last.visit, "%m/%d/%y")
haart$date.death <- as.Date(haart$date.death, "%m/%d/%y")
haart$indicator <- ifelse((haart$date.death - haart$init.date > 365 | is.na(haart$date.death)), 0, 1)
haart$follow.up <- ifelse(is.na(haart$date.death), haart$last.visit - haart$init.date,
                          haart$date.death - haart$init.date)
haart$follow.up[haart$follow.up > 365] <- 365
haart$loss.to.followup <- ifelse(haart$follow.up < 365 & haart$death == 0, 1, 0)
# create a list for all users
row.reg <- strsplit(haart[, 'init.reg'], ",")
# find all the unique drugs
all.reg <- unique(unlist(strsplit(haart[, 'init.reg'], ",")))
# create a database for user.reg
```

```

user.reg <- sapply(all.reg, function(j) sapply(row.reg, function(i) j %in% i))
# append it to haart database
haart <- cbind(haart, user.reg)
head(haart, 5)

```

```

##   male age aids cd4baseline logvl  weight hemoglobin  init.reg
## 1    1  25    0          NA    NA      NA          NA 3TC,AZT,EFV
## 2    1  49    0         143    NA  58.0608         11 3TC,AZT,EFV
## 3    1  42    1         102    NA  48.0816          1 3TC,AZT,EFV
## 4    0  33    0         107    NA  46.0000         NA 3TC,AZT,NVP
## 5    1  27    0          52     4     NA          NA 3TC,D4T,EFV
##   init.date last.visit death date.death indicator follow.up
## 1 2003-07-01 2007-02-26     0      <NA>         0       365
## 2 2004-11-23 2008-02-22     0      <NA>         0       365
## 3 2003-04-30 2005-11-21     1 2006-01-11         0       365
## 4 2006-03-25 2006-05-05     1 2006-05-07         1        43
## 5 2004-09-01 2007-11-13     0      <NA>         0       365
##   loss.to.followup 3TC  AZT  EFV  NVP  D4T  ABC  DDI  IDV  LPV
## 1                  0 TRUE  TRUE  TRUE FALSE FALSE FALSE FALSE FALSE
## 2                  0 TRUE  TRUE  TRUE FALSE FALSE FALSE FALSE FALSE
## 3                  0 TRUE  TRUE  TRUE FALSE FALSE FALSE FALSE FALSE
## 4                  0 TRUE  TRUE FALSE  TRUE FALSE FALSE FALSE FALSE
## 5                  0 TRUE FALSE  TRUE FALSE  TRUE FALSE FALSE FALSE
##   RTV  SQV  FTC  TDF  DDC  NFV  T20  ATV  FPV
## 1 FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## 2 FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## 3 FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## 4 FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## 5 FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE

```

```
tail(haart, 5)
```

```

##   male      age aids cd4baseline  logvl  weight hemoglobin
## 1000    0 40.00000    1        131    NA  46.2672          8
## 1001    0 27.00000    0        232    NA     NA          NA
## 1002    1 38.72142    0        170    NA  84.0000          NA
## 1003    1 23.00000   NA        154  3.995635  65.5000         14
## 1004    0 31.00000    0        236    NA  45.8136          NA
##   init.reg  init.date last.visit death date.death indicator
## 1000 3TC,D4T,NVP 2003-07-03 2008-02-29     0      <NA>         0
## 1001 3TC,AZT,NVP 2003-12-01 2004-01-05     0      <NA>         0
## 1002 3TC,AZT,NVP 2002-09-26 2004-03-29     0      <NA>         0
## 1003 3TC,DDI,EFV 2007-01-31 2007-04-16     0      <NA>         0
## 1004 3TC,D4T,NVP 2003-12-03 2007-10-11     0      <NA>         0
##   follow.up loss.to.followup 3TC  AZT  EFV  NVP  D4T  ABC  DDI
## 1000       365                0 TRUE FALSE TRUE  TRUE TRUE FALSE
## 1001        35                1 TRUE TRUE FALSE TRUE FALSE FALSE
## 1002       365                0 TRUE TRUE FALSE TRUE FALSE FALSE
## 1003        75                1 TRUE FALSE TRUE FALSE FALSE TRUE
## 1004       365                0 TRUE FALSE FALSE TRUE TRUE FALSE
##   IDV  LPV  RTV  SQV  FTC  TDF  DDC  NFV  T20  ATV  FPV
## 1000 FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## 1001 FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## 1002 FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE

```

```
## 1003 FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## 1004 FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
```

Question 2

14 points

Use the following code to generate data for patients with repeated measures of A1C (a test for levels of blood glucose).

```
genData <- function(n) {
  if(exists(".Random.seed", envir = .GlobalEnv)) {
    save.seed <- get(".Random.seed", envir = .GlobalEnv)
    on.exit(assign(".Random.seed", save.seed, envir = .GlobalEnv))
  } else {
    on.exit(rm(".Random.seed", envir = .GlobalEnv))
  }
  set.seed(n)
  subj <- ceiling(n / 10)
  id <- sample(subj, n, replace=TRUE)
  times <- as.integer(difftime(as.POSIXct("2005-01-01"), as.POSIXct("2000-01-01"), units='secs'))
  dt <- as.POSIXct(sample(times, n), origin='2000-01-01')
  mu <- runif(subj, 4, 10)
  a1c <- unsplit(mapply(rnorm, tabulate(id), mu, SIMPLIFY=FALSE), id)
  data.frame(id, dt, a1c)
}
x <- genData(500)
```

Perform the following manipulations: (2 points each)

1. Order the data set by id and dt.

```
x <- x[order(x$id, x$dt), ]
row.names(x) <- NULL
```

2. For each id, determine if there is more than a one year gap in between observations. Add a new row at the one year mark, with the a1c value set to missing. A two year gap would require two new rows, and so forth.

```
library("lubridate")
```

```
##
## Attaching package: 'lubridate'
## The following object is masked from 'package:base':
##
##      date
x_new <- x
# append the new row after x_new
for (i in seq(nrow(x)-1)) {
  # if the same id
  if (x$id[i]==x$id[i+1]) {
    temp <- floor(difftime(x$dt[i+1], x$dt[i], units='days')/365)
  }
  if (temp > 0) {
    # create a row-size(temp) dataframe as newrow
    newrow <- data.frame(matrix(nrow = temp, ncol = 3))
```

```

newrow[, 1] <- x$id[i]
newrow[, 2] <- as.character.Date(x$dt[i] + years(1:temp))
colnames(newrow) <- c("id", "dt", "a1c")
x_new <- rbind(x_new, newrow)
}
}
# set x_new to x back
x <- x_new
# reorder x by id and dt
x <- x[order(x$id, x$dt), ]
row.names(x) <- NULL

```

3. Create a new column visit. For each id, add the visit number. This should be 1 to n where n is the number of observations for an individual. This should include the observations created with missing a1c values.

```

x$visit <- 1
id_temp <- x$id[1]
for (i in 2:nrow(x)) {
  # if found not the same id, change id_temp
  if (x$id[i] != id_temp) {id_temp <- x$id[i]}
  else {
    x$visit[i] <- x$visit[i-1] + 1
  }
}

```

4. For each id, replace missing values with the mean a1c value for that individual.

```

for (i in 1:nrow(x)) {
  if(is.na(x$a1c[i])) {
    x$a1c[i] <- mean(x$a1c[which(x$id == x$id[i])], na.rm = TRUE)
  }
}

```

5. Print mean a1c for each id.

```

id_list <- unique(x$id)
for (i in id_list) {
  print(i)
  print(mean(x$a1c[x$id == i]))
}

```

```

## [1] 1
## [1] 4.063372
## [1] 2
## [1] 7.544643
## [1] 3
## [1] 6.75764
## [1] 4
## [1] 3.892127
## [1] 5
## [1] 9.512311
## [1] 6
## [1] 7.555965
## [1] 7
## [1] 9.161686

```

```
## [1] 8
## [1] 7.189064
## [1] 9
## [1] 9.283873
## [1] 10
## [1] 7.975217
## [1] 11
## [1] 6.917562
## [1] 12
## [1] 7.034021
## [1] 13
## [1] 9.145282
## [1] 14
## [1] 6.623756
## [1] 15
## [1] 8.012406
## [1] 16
## [1] 4.222158
## [1] 17
## [1] 3.996034
## [1] 18
## [1] 9.164873
## [1] 19
## [1] 5.50721
## [1] 20
## [1] 3.726675
## [1] 21
## [1] 8.140939
## [1] 22
## [1] 5.637501
## [1] 23
## [1] 7.366889
## [1] 24
## [1] 7.439316
## [1] 25
## [1] 6.877135
## [1] 26
## [1] 6.556759
## [1] 27
## [1] 4.926457
## [1] 28
## [1] 7.433917
## [1] 29
## [1] 4.508086
## [1] 30
## [1] 6.045577
## [1] 31
## [1] 7.116586
## [1] 32
## [1] 6.568791
## [1] 33
## [1] 6.494069
## [1] 34
## [1] 6.768615
```

```
## [1] 35
## [1] 8.4767
## [1] 36
## [1] 9.60441
## [1] 37
## [1] 9.606253
## [1] 38
## [1] 5.355979
## [1] 39
## [1] 6.917013
## [1] 40
## [1] 9.530136
## [1] 41
## [1] 9.802424
## [1] 42
## [1] 3.89177
## [1] 43
## [1] 6.095849
## [1] 44
## [1] 9.09167
## [1] 45
## [1] 6.737204
## [1] 46
## [1] 9.621763
## [1] 47
## [1] 9.231489
## [1] 48
## [1] 6.4046
## [1] 49
## [1] 6.096076
## [1] 50
## [1] 8.962319
```

6. Print total number of visits for each id.

```
for (i in id_list) {
  print(max(x$visit[x$id == i]))
}
```

```
## [1] 11
## [1] 20
## [1] 14
## [1] 12
## [1] 14
## [1] 10
## [1] 10
## [1] 12
## [1] 11
## [1] 12
## [1] 10
## [1] 10
## [1] 9
## [1] 12
## [1] 9
## [1] 11
```

```
## [1] 12
## [1] 10
## [1] 11
## [1] 9
## [1] 10
## [1] 8
## [1] 8
## [1] 15
## [1] 12
## [1] 14
## [1] 11
## [1] 14
## [1] 10
## [1] 7
## [1] 11
## [1] 5
## [1] 9
## [1] 12
## [1] 11
## [1] 9
## [1] 17
## [1] 15
## [1] 8
## [1] 8
## [1] 17
## [1] 14
## [1] 11
## [1] 12
## [1] 14
## [1] 9
## [1] 12
## [1] 11
## [1] 12
## [1] 10
```

7. Print the observations for id = 15.

```
print(x[x$id == 15, ])
```

```
##      id      dt      a1c visit
## 168 15 2000-04-30 00:34:50 7.527105      1
## 169 15 2001-01-17 21:11:02 5.898371      2
## 170 15 2001-04-25 06:23:05 8.566593      3
## 171 15 2002-04-25 06:23:05 8.012406      4
## 172 15 2003-04-25 06:23:05 8.012406      5
## 173 15 2003-06-06 14:06:00 9.133769      6
## 174 15 2004-06-06 14:06:00 8.012406      7
## 175 15 2004-08-20 17:47:11 8.936190      8
## 176 15 2005-08-20 17:47:11 8.012406      9
```

Question 3

10 points

Import the `addr.txt` file from the GitHub repository. This file contains a listing of names and addresses

(thanks google). Parse each line to create a data.frame with the following columns: lastname, firstname, streetno, streetname, city, state, zip. Keep middle initials or abbreviated names in the firstname column. Print out the entire data.frame.

```
url <- "https://github.com/fonnesbeck/Bios6301/raw/master/datasets/addr.txt"
addr <- read.table(url, header = FALSE, sep="\t", colClasses=c("character"))
# create a list
temp <- unlist(strsplit(addr[,1], " "))

# remove the " " on two sides in each element, replaced by ""
trim <- function (x) gsub("^\\s+|\\s+$", "", x)
temp <- trim(temp)

# remove the "" in the list
temp <- temp[temp != ""]
# create a matrix by each 6 elements by first, last, addr, city, state, zip
dat <- matrix(temp, ncol = 6, byrow = TRUE)

# need to split addr into streetno and streetname
# we observe the streetno, 1 space, streetname
# the expression starts with (words), 1 space and ends with any character 0 or any times
y <- data.frame(streetno=sub("^([0-9]{3,4})\\s?(.*)$", "\\1", dat[,3]), streetname=sub("^([0-9]{3,4})\\s?", "", dat[,3]))
dat <- dat[,-3]
dat <- cbind(y, dat)
colnames(dat) <- c("streetno", "streetname", "lastname", "firstname", "city", "state", "zip")
# reorder the dataframe
dat_new <- dat[, c(3,4,1,2,5,6,7)]
print(dat_new)
```

##	lastname	firstname	streetno	streetname	city	state
## 1	Bania	Thomas M.	725	Commonwealth Ave.	Boston	MA
## 2	Barnaby	David	373	W. Geneva St.	Wms. Bay	WI
## 3	Bausch	Judy	373	W. Geneva St.	Wms. Bay	WI
## 4	Bolatto	Alberto	725	Commonwealth Ave.	Boston	MA
## 5	Carlstrom	John	933	E. 56th St.	Chicago	IL
## 6	Chamberlin	Richard A.	111	Nowelo St.	Hilo	HI
## 7	Chuss	Dave	2145	Sheridan Rd	Evanston	IL
## 8	Davis	E. J.	933	E. 56th St.	Chicago	IL
## 9	Depoy	Darren	174	W. 18th Ave.	Columbus	OH
## 10	Griffin	Greg	5000	Forbes Ave.	Pittsburgh	PA
## 11	Halvorsen	Nils	933	E. 56th St.	Chicago	IL
## 12	Harper	Al	373	W. Geneva St.	Wms. Bay	WI
## 13	Huang	Maohai	725	W. Commonwealth Ave.	Boston	MA
## 14	Ingalls	James G.	725	W. Commonwealth Ave.	Boston	MA
## 15	Jackson	James M.	725	W. Commonwealth Ave.	Boston	MA
## 16	Knudsen	Scott	373	W. Geneva St.	Wms. Bay	WI
## 17	Kovac	John	5640	S. Ellis Ave.	Chicago	IL
## 18	Landsberg	Randy	5640	S. Ellis Ave.	Chicago	IL
## 19	Lo	Kwok-Yung	1002	W. Green St.	Urbana	IL
## 20	Loewenstein	Robert F.	373	W. Geneva St.	Wms. Bay	WI
## 21	Lynch	John	4201	Wilson Blvd	Arlington	VA
## 22	Martini	Paul	174	W. 18th Ave.	Columbus	OH
## 23	Meyer	Stephan	933	E. 56th St.	Chicago	IL
## 24	Mrozek	Fred	373	W. Geneva St.	Wms. Bay	WI
## 25	Newcomb	Matt	5000	Forbes Ave.	Pittsburgh	PA

## 26	Novak	Giles	2145	Sheridan Rd	Evanston	IL
## 27	Odalen	Nancy	373	W. Geneva St.	Wms. Bay	WI
## 28	Pernic	Dave	373	W. Geneva St.	Wms. Bay	WI
## 29	Pernic	Bob	373	W. Geneva St.	Wms. Bay	WI
## 30	Peterson	Jeffrey	5000	Forbes Ave.	Pittsburgh	PA
## 31	Pryke	Clem	933	E. 56th St.	Chicago	IL
## 32	Rebull	Luisa	5640	S. Ellis Ave.	Chicago	IL
## 33	Renbarger	Thomas	2145	Sheridan Rd	Evanston	IL
## 34	Rottman	Joe	8730	W. Mountain View Ln	Littleton	CO
## 35	Schartman	Ethan	933	E. 56th St.	Chicago	IL
## 36	Spotz	Bob	373	W. Geneva St.	Wms. Bay	WI
## 37	Thoma	Mark	373	W. Geneva St.	Wms. Bay	WI
## 38	Walker	Chris	933	N. Cherry St.	Tucson	AZ
## 39	Wehrer	Cheryl	5000	Forbes Ave.	Pittsburgh	PA
## 40	Wirth	Jesse	373	W. Geneva St.	Wms. Bay	WI
## 41	Wright	Greg	791	Holmdel-Keyport Rd.	Holmdel	NY
## 42	Zingale	Michael	5640	S. Ellis Ave.	Chicago	IL
##	zip					
## 1	02215					
## 2	53191					
## 3	53191					
## 4	02215					
## 5	60637					
## 6	96720					
## 7	60208-3112					
## 8	60637					
## 9	43210					
## 10	15213					
## 11	60637					
## 12	53191					
## 13	02215					
## 14	02215					
## 15	02215					
## 16	53191					
## 17	60637					
## 18	60637					
## 19	61801					
## 20	53191					
## 21	22230					
## 22	43210					
## 23	60637					
## 24	53191					
## 25	15213					
## 26	60208-3112					
## 27	53191					
## 28	53191					
## 29	53191					
## 30	15213					
## 31	60637					
## 32	60637					
## 33	60208-3112					
## 34	80125					
## 35	60637					
## 36	53191					

```
## 37      53191
## 38      85721
## 39      15213
## 40      53191
## 41 07733-1988
## 42      60637
```

Question 4

2 points

The first argument to most functions that fit linear models are formulas. The following example defines the response variable `death` and allows the model to incorporate all other variables as terms. `.` is used to mean all columns not otherwise in the formula.

```
url <- "https://github.com/fonnesbeck/Bios6301/raw/master/datasets/haart.csv"
haart_df <- read.csv(url)[,c('death', 'weight', 'hemoglobin', 'cd4baseline')]
coef(summary(glm(death ~ ., data=haart_df, family=binomial(logit))))
```

```
##              Estimate Std. Error  z value    Pr(>|z|)
## (Intercept)  3.576411744 1.226870535  2.915069 0.0035561039
## weight      -0.046210552 0.022556001 -2.048703 0.0404911395
## hemoglobin   -0.350642786 0.105064078 -3.337418 0.0008456055
## cd4baseline  0.002092582 0.001811959  1.154872 0.2481427160
```

Now imagine running the above several times, but with a different response and data set each time. Here's a function:

```
myfun <- function(dat, response) {
  form <- as.formula(response ~ .)
  coef(summary(glm(form, data=dat, family=binomial(logit))))
}
```

Unfortunately, it doesn't work. `tryCatch` is "catching" the error so that this file can be knit to PDF.

```
tryCatch(myfun(haart_df, death), error = function(e) e)
```

```
## <simpleError in eval(expr, envir, enclos): object 'death' not found>
```

What do you think is going on? Consider using `debug` to trace the problem.

`debug(myfun(haart_df, death))` can not find death debug at #2: `form <- as.formula(response ~ .)` The bug: `as.formula` function needs a character object variable, use `paste` function. (check `as.formula` function in the help for details)

5 bonus points

Create a working function.

```
myfun <- function(dat, response) {
  # make ~ . as character
  form <- as.formula(paste(response, "~ ."))
  coef(summary(glm(form, data=dat, family=binomial(logit))))
}
# use death as character
myfun(haart_df, 'death')
```

```
##              Estimate Std. Error  z value    Pr(>|z|)
## (Intercept)  3.576411744 1.226870535  2.915069 0.0035561039
```

```
## weight      -0.046210552 0.022556001 -2.048703 0.0404911395
## hemoglobin  -0.350642786 0.105064078 -3.337418 0.0008456055
## cd4baseline 0.002092582 0.001811959 1.154872 0.2481427160
```